

```
import torch
import torch.nn as nn
import torch.nn.functional as F
import torchvision.transforms as transforms
import torch.optim as optim
import torchvision
from torch.utils.data import DataLoader, Dataset
from skimage import io, color
import pandas as pd
import os
from tqdm.notebook import tqdm
```

```
# Takes inputs with dims = (N, C, *)
# Gives outputs with dimes = (N, C, *)
class LocalResponseNormalization(nn.Module):
    def __init__(self, neighbourhood_length, normalisation_const_alpha, contrast_const_beta, noise_k):
        super(LocalResponseNormalization, self).__init__()
        self.nbd_len = neighbourhood_length
        self.alpha = normalisation_const_alpha
        self.beta = contrast_const_beta
        self.k = noise_k

    # The following is exactly what pytorch does under the hood as well. I only replicated it for my understanding :)
    def forward(self, x):
        # Lets validate if x is atleast 3 dimensional
        dim = x.dim()
        if dim < 3:
            raise ValueError("Expected tensor of atleast 3 dimensions, found only {}".format(dim))
        denom = x.pow(2).unsqueeze(1)
        if dim == 3:
            denom = F.pad(denom, (0, 0, self.nbd_len // 2, (self.nbd_len - 1) // 2))
            denom = F.avg_pool2d(denom, (self.nbd_len, 1), stride=1)
            denom = denom.squeeze(1)
        else:
            sizes = x.size()
            # The last two dimensions make up a single channel. The third dimension decides the number of channels
            # across which we will apply local response normalization.
            denom = denom.view(sizes[0], 1, sizes[1], sizes[2], -1)
            # The point is to pad in front and back of the channels across which we'll apply normalization
            denom = F.pad(denom, (0, 0, 0, 0, self.nbd_len // 2, (self.nbd_len - 1) // 2))
            denom = F.avg_pool3d(denom, (self.nbd_len, 1, 1), stride=1)
            denom = denom.squeeze(1).view(sizes)
        denom = denom.mul(self.alpha).add(self.k).pow(self.beta)
        return x.div(denom)
```

```
# Expects input tensor to be of dimensions (batch_size, 3, 224, 224)
class Alexnet(nn.Module):
    def __init__(self):
        super(Alexnet, self).__init__()
        self.conv1 = nn.Conv2d(in_channels=3, out_channels=96, kernel_size=11, stride=4, padding=2)
        self.conv2 = nn.Conv2d(in_channels=96, out_channels=256, kernel_size=5, stride=1, padding=2)
        self.conv3 = nn.Conv2d(in_channels=256, out_channels=384, kernel_size=3, stride=1, padding=1)
        self.conv4 = nn.Conv2d(in_channels=384, out_channels=384, kernel_size=3, stride=1, padding=1)
        self.conv5 = nn.Conv2d(in_channels=384, out_channels=256, kernel_size=3, stride=1, padding=1)
        self.fc1 = nn.Linear(in_features=256 * 6 * 6, out_features=4096)
        self.fc2 = nn.Linear(in_features=4096, out_features=4096)
        self.fc3 = nn.Linear(in_features=4096, out_features=10)
        self.max_pool = nn.MaxPool2d(kernel_size=3, stride=2)
        # This layer helps us avoid calculating output map size when feeding into a linear layer in PyTorch.
        self.adaptive_pool = nn.AdaptiveAvgPool2d(output_size=(6, 6))
        self.norm = LocalResponseNormalization(neighbourhood_length=5, normalisation_const_alpha=1e-4, contrast_const_beta=0.75, noise_k=
self.dropout = nn.Dropout())

    def forward(self, x):
        x = self.max_pool(self.norm(F.relu(self.conv1(x))))
        x = self.max_pool(self.norm(F.relu(self.conv2(x))))
        x = F.relu(self.conv3(x))
        x = F.relu(self.conv4(x))
        x = self.adaptive_pool(self.norm(F.relu(self.conv5(x))))
        x = torch.flatten(x, 1)
        x = F.relu(self.fc1(x))
        x = F.relu(self.fc2(x))
        x = self.dropout(x)
        x = self.fc3(x)
        return x
```

```
class SceneDataset(Dataset):
    def __init__(self, annotations_csv, root_dir, transform=None):
        self.annotations = pd.read_csv(annotations_csv)
```

```

        self.root_dir = root_dir
        self.transform = transform

    def __len__(self):
        return len(self.annotations)

    def __getitem__(self, index):
        img_path = os.path.join(self.root_dir, self.annotations.iloc[index, 0])
        image = io.imread(img_path)
        label = torch.tensor(int(self.annotations.iloc[index, 1]))
        if self.transform:
            image = self.transform(image)
        return [image, label]

```

```

def check_accuracy(loader, model):
    num_correct = 0
    num_samples = 0
    # Don't forget to toggle to eval mode!
    model.eval()

    with torch.no_grad():
        for data, targets in tqdm(loader):
            data = data.to(device)
            targets = targets.to(device)
            scores = model(data)
            _, predictions = scores.max(1)
            num_correct += (predictions == targets).sum()
            num_samples += predictions.size(0)

    print("Correct: {}, Total: {}, Accuracy: {}".format(num_correct, num_samples, int(num_correct) / int(num_samples)))
    # Don't forget to toggle back to model.train() since you're done with evaluation
    model.train()

```

```

from google.colab import drive
drive.mount('/content/drive')

```

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True).

```
%cd /content/drive/My Drive/KIST Hackathon 2023/Phase II workshop/Day 1/files and datasets/files
```

```
/content/drive/My Drive/KIST Hackathon 2023/Phase II workshop/Day 1/files and datasets/files
```

```

if __name__ == '__main__':

    LEARNING_RATE = 0.0001
    # You could try playing around with the batch size(say 16) and learning rate(say 0.001) for faster convergence.
    BATCH_SIZE = 8
    EPOCHS = 10
    device = "cuda" if torch.cuda.is_available() else "cpu"

    transform_img = transforms.Compose([
        transforms.ToPILImage(),
        transforms.Resize((224, 224)),
        transforms.ToTensor()
    ])

    data = SceneDataset(annotations_csv="train-scene classification/train.csv",
                        root_dir="train-scene classification/train",
                        transform = transform_img)

    print(len(data))
    train_data, test_data = torch.utils.data.random_split(data, [65, 9])
    train_loader = DataLoader(dataset=train_data, batch_size=BATCH_SIZE, shuffle=True)
    test_loader = DataLoader(dataset=test_data, batch_size=BATCH_SIZE, shuffle=True)

    alexnet = Alexnet()
    alexnet.to(device)
    criterion = nn.CrossEntropyLoss()
    optimizer = optim.Adam(alexnet.parameters(), lr=LEARNING_RATE)

    data, targets = next(iter(train_loader))
    for epoch in tqdm(range(EPOCHS)):
        losses = []
        with tqdm(total=len(train_loader)) as pbar:
            for batch_idx, (data, targets) in enumerate(train_loader):
                data = data.to(device=device)
                targets = targets.to(device=device)

                scores = alexnet(data)
                loss = criterion(scores, targets)
                losses.append(loss)

```

```
        # backprop
        optimizer.zero_grad()
        loss.backward()
        optimizer.step()
#         print(loss.item())
        pbar.update(1)
    print("Cost at epoch {} is {}".format(epoch, sum(losses) / len(losses)))
    check_accuracy(train_loader, alexnet)
    check_accuracy(test_loader, alexnet)
```



```
Correct: 2, Total: 9, Accuracy: 0.2222222222222222
100% 2/2 [00:00<00:00, 2.54it/s]

Correct: 2, Total: 9, Accuracy: 0.2222222222222222
100% 9/9 [00:13<00:00, 1.20s/it]

Cost at epoch 2 is 1.7962262630462646
100% 9/9 [00:03<00:00, 2.34it/s]

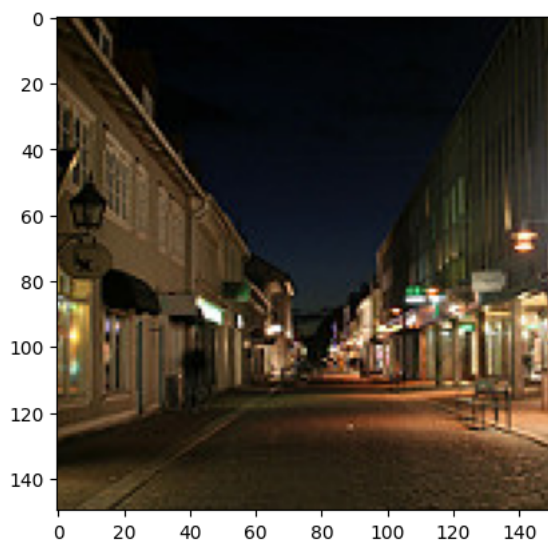
Correct: 20, Total: 65, Accuracy: 0.3076923076923077
100% 2/2 [00:00<00:00, 2.57it/s]

Correct: 1, Total: 9, Accuracy: 0.1111111111111111
```

```
import matplotlib.pyplot as plt
```

```
def show_image_and_prediction(img_path):
    idx2label = {0: "Buildings", 1: "Forests", 2: "Mountains", 3: "Glacier", 4: "Sea", 5: "Street"}
    img = io.imread(img_path)
    transformed_img = transform_img(img)
    out = alexnet(transformed_img.unsqueeze(0).to(device=device))
    _, pred = out.max(1)
    plt.imshow(img)
    plt.show()
    print("Alexnet thinks this scene contains: {}".format(idx2label[pred.item()]))
```

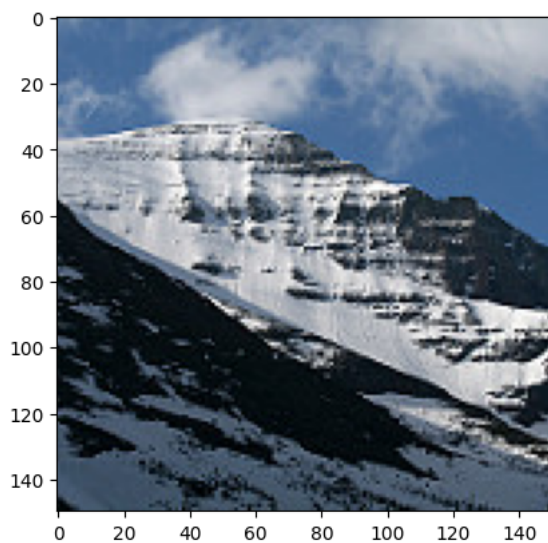
```
show_image_and_prediction("train-scene classification/train/50.jpg")
```



Alexnet thinks this scene contains: Street

```
Correct: 1, Total: 9, Accuracy: 0.1111111111111111
```

```
show_image_and_prediction("train-scene classification/train/30.jpg")
```



Alexnet thinks this scene contains: Street

Yolo Model:

```
!pip install ultralytics==8.0.20
```


```
from IPython import display
display.clear_output()

from ultralytics import YOLO
from IPython.display import display, Image
```

```
!yolo task=detect mode=predict model=yolov8n.pt conf=0.25 source='images2.jpg' save=true
```

2023-09-12 14:09:03.585009: I tensorflow/core/platform/cpu_feature_guard.cc:182] This TensorFlow binary is optimized to use available AVX2 FMA, in other operations, rebuild TensorFlow with the appropriate compiler flags.
 2023-09-12 14:09:04.751814: W tensorflow/compiler/tf2tensorrt/utils/py_utils.cc:38] TF-TRT Warning: Could not find TensorRT
 Ultralytics YOLOv8.0.20 Python-3.10.12 torch-2.0.1+cu118 CPU
 YOLOv8n summary (fused): 168 layers, 3151904 parameters, 0 gradients, 8.7 GFLOPs
 image 1/1 /content/drive/MyDrive/KIST Hackathon 2023/Phase II workshop/Day 1/files and datasets/files/images2.jpg: 448x640 2 zebras
 Speed: 1.3ms pre-process, 115.1ms inference, 2.9ms postprocess per image at shape (1, 3, 640, 640)
 Results saved to runs/detect/predict3

```
Image(filename='runs/detect/predict3/images2.jpg', height=500)
```



Yolo live detection

```
from IPython.display import display, Javascript, Image
from google.colab.output import eval_js
from google.colab.patches import cv2_imshow
from base64 import b64decode, b64encode
import cv2
import numpy as np
import PIL
import io
import html
import time
import matplotlib.pyplot as plt
%matplotlib inline
```

```
# clone darknet repo
!git clone https://github.com/AlexeyAB/darknet
```

```
Cloning into 'darknet'...
remote: Enumerating objects: 15750, done.
remote: Counting objects: 100% (213/213), done.
remote: Compressing objects: 100% (125/125), done.
remote: Total 15750 (delta 103), reused 155 (delta 86), pack-reused 15537
Receiving objects: 100% (15750/15750), 14.62 MiB | 11.82 MiB/s, done.
Resolving deltas: 100% (10536/10536), done.
Updating files: 100% (2063/2063), done.
```

```
# change makefile to have GPU, OPENCV and LIBSO enabled
%cd darknet
!sed -i 's/OPENCV=0/OPENCV=1/' Makefile
!sed -i 's/GPU=0/GPU=1/' Makefile
!sed -i 's/CUDNN=0/CUDNN=1/' Makefile
```

```
!sed -i 's/CUDNN_HALF=0/CUDNN_HALF=1/' Makefile
!sed -i 's/LIBSO=0/LIBSO=1/' Makefile
```

/content/drive/MyDrive/KIST Hackathon 2023/Phase II workshop/Day 1/files and datasets/files/darknet

!make

```
mkdir -p ./obj/
mkdir -p backup
chmod +x *.sh
g++ -std=c++11 -std=c++11 -Iinclude/ -I3rdparty/stb/include -DOPENCV `pkg-config --cflags opencv4 2> /dev/null || pkg-config --c
./src/image_opencv.cpp: In function 'void draw_detections_cv_v3(void**, detection*, int, float, char**, image**, int, int)':
./src/image_opencv.cpp:946:23: warning: variable 'rgb' set but not used [-Wunused-but-set-variable]
  946 |         float rgb[3];
      |         ^~~~
./src/image_opencv.cpp: In function 'void cv_draw_object(image, float*, int, int, int*, float*, int*, int, char**)':
./src/image_opencv.cpp:1444:14: warning: unused variable 'buff' [-Wunused-variable]
 1444 |         char buff[100];
      |         ^~~~~
./src/image_opencv.cpp:1420:9: warning: unused variable 'it_tb_res' [-Wunused-variable]
 1420 |         int it_tb_res = cv::createTrackbar(it_trackbar_name, window_name, &it_trackbar_value, 1000);
      |         ^~~~~~
./src/image_opencv.cpp:1424:9: warning: unused variable 'lr_tb_res' [-Wunused-variable]
 1424 |         int lr_tb_res = cv::createTrackbar(lr_trackbar_name, window_name, &lr_trackbar_value, 20);
      |         ^~~~~~
./src/image_opencv.cpp:1428:9: warning: unused variable 'cl_tb_res' [-Wunused-variable]
 1428 |         int cl_tb_res = cv::createTrackbar(cl_trackbar_name, window_name, &cl_trackbar_value, classes-1);
      |         ^~~~~~
./src/image_opencv.cpp:1431:9: warning: unused variable 'bo_tb_res' [-Wunused-variable]
 1431 |         int bo_tb_res = cv::createTrackbar(bo_trackbar_name, window_name, boxonly, 1);
      |         ^~~~~~
g++ -std=c++11 -std=c++11 -Iinclude/ -I3rdparty/stb/include -DOPENCV `pkg-config --cflags opencv4 2> /dev/null || pkg-config --c
./src/http_stream.cpp: In member function 'bool JSON_sender::write(const char*)':
./src/http_stream.cpp:253:21: warning: unused variable 'n' [-Wunused-variable]
  253 |         int n = _write(client, outputbuf, outlen);
      |         ^
./src/http_stream.cpp: In function 'void set_track_id(detection*, int, float, float, float, int, int, int)':
./src/http_stream.cpp:867:27: warning: comparison of integer expressions of different signedness: 'int' and 'std::vector<detection*
  867 |         for (int i = 0; i < v.size(); ++i) {
      |         ~~~~~
./src/http_stream.cpp:875:33: warning: comparison of integer expressions of different signedness: 'int' and 'std::vector<detection*
  875 |         for (int old_id = 0; old_id < old_dets.size(); ++old_id) {
      |         ~~~~~
./src/http_stream.cpp:894:31: warning: comparison of integer expressions of different signedness: 'int' and 'std::vector<detection*
  894 |         for (int index = 0; index < new_dets_num*old_dets.size(); ++index) {
      |         ~~~~~
./src/http_stream.cpp:930:28: warning: comparison of integer expressions of different signedness: 'std::deque<std::vector<detection*
  930 |         if (old_dets_dq.size() > deque_size) old_dets_dq.pop_front();
      |         ~~~~~
gcc -Iinclude/ -I3rdparty/stb/include -DOPENCV `pkg-config --cflags opencv4 2> /dev/null || pkg-config --cflags opencv` -DGPU -I
./src/gemm.c: In function 'convolution_2d':
./src/gemm.c:2044:15: warning: unused variable 'out_w' [-Wunused-variable]
 2044 |         const int out_w = (w + 2 * pad - ksize) / stride + 1; // output_width=input_width for stride=1 and pad=1
      |         ^~~~~
./src/gemm.c:2043:15: warning: unused variable 'out_h' [-Wunused-variable]
 2043 |         const int out_h = (h + 2 * pad - ksize) / stride + 1; // output_height=input_height for stride=1 and pad=1
      |         ^~~~~
gcc -Iinclude/ -I3rdparty/stb/include -DOPENCV `pkg-config --cflags opencv4 2> /dev/null || pkg-config --cflags opencv` -DGPU -I
./src/utils.c: In function 'custom_hash':
./src/utils.c:1093:12: warning: suggest parentheses around assignment used as truth value [-Wparentheses]
 1093 |         while (c = *str++)
      |         ^
In file included from /usr/include/string.h:535,
      from include/darknet.h:14,
```

```
# get bthe scaled yolov4 weights file that is pre-trained to detect 80 classes (objects) from shared google drive
!wget --load-cookies /tmp/cookies.txt "https://docs.google.com/uc?export=download&confirm=$(wget --quiet --save-cookies /tmp/cookies.txt
```

```
--2023-09-12 14:22:03-- https://docs.google.com/uc?export=download&confirm=t&id=1V3vsIaxAlGwV4Aar9bAik5U0QFttKwg
Resolving docs.google.com (docs.google.com)... 74.125.142.102, 74.125.142.101, 74.125.142.139, ...
Connecting to docs.google.com (docs.google.com)|74.125.142.102|:443... connected.
HTTP request sent, awaiting response... 303 See Other
Location: https://doc-14-84-docs.googleusercontent.com/docs/securesc/ha0ro937gcuc7l7deffksulhg5h7mbp1/1ilu0v3i7v7on09ejrqa2ch7seoun
Warning: wildcards not supported in HTTP.
--2023-09-12 14:22:04-- https://doc-14-84-docs.googleusercontent.com/docs/securesc/ha0ro937gcuc7l7deffksulhg5h7mbp1/1ilu0v3i7v7on0
Resolving doc-14-84-docs.googleusercontent.com (doc-14-84-docs.googleusercontent.com)... 74.125.197.132, 2607:f8b0:400e:c03::84
Connecting to doc-14-84-docs.googleusercontent.com (doc-14-84-docs.googleusercontent.com)|74.125.197.132|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 211944840 (202M) [application/octet-stream]
Saving to: 'yolov4-csp.weights'

yolov4-csp.weights 100%[=====] 202.13M 44.9MB/s in 4.5s

2023-09-12 14:22:08 (44.6 MB/s) - 'yolov4-csp.weights' saved [211944840/211944840]
```

```
# import darknet functions to perform object detections
from darknet import *
# load in our YOLOv4 architecture network
network, class_names, class_colors = load_network("cfg/yolov4-csp.cfg", "cfg/coco.data", "yolov4-csp.weights")
width = network_width(network)
height = network_height(network)

# darknet helper function to run detection on image
def darknet_helper(img, width, height):
    darknet_image = make_image(width, height, 3)
    img_rgb = cv2.cvtColor(img, cv2.COLOR_BGR2RGB)
    img_resized = cv2.resize(img_rgb, (width, height),
                              interpolation=cv2.INTER_LINEAR)

    # get image ratios to convert bounding boxes to proper size
    img_height, img_width, _ = img.shape
    width_ratio = img_width/width
    height_ratio = img_height/height

    # run model on darknet style image to get detections
    copy_image_from_bytes(darknet_image, img_resized.tobytes())
    detections = detect_image(network, class_names, darknet_image)
    free_image(darknet_image)
    return detections, width_ratio, height_ratio
```

```
-----
OSError                                Traceback (most recent call last)
<ipython-input-63-5a31be365992> in <cell line: 2>()
      1 # import darknet functions to perform object detections
----> 2 from darknet import *
      3 # load in our YOLOv4 architecture network
      4 network, class_names, class_colors = load_network("cfg/yolov4-csp.cfg", "cfg/coco.data",
"yolov4-csp.weights")
      5 width = network_width(network)
```

```
----- 1 frames -----
/usr/lib/python3.10/ctypes/__init__.py in __init__(self, name, mode, handle, use_errno,
use_last_error, winmode)
    372
    373     if handle is None:
--> 374         self._handle = _dlopen(self._name, mode)
    375     else:
    376         self._handle = handle
```

```
OSError: libcuda.so.1: cannot open shared object file: No such file or directory
```

```
# run test on person.jpg image that comes with repository
image = cv2.imread("images1.jpeg")
detections, width_ratio, height_ratio = darknet_helper(image, width, height)

for label, confidence, bbox in detections:
    left, top, right, bottom = bbox2points(bbox)
    left, top, right, bottom = int(left * width_ratio), int(top * height_ratio), int(right * width_ratio), int(bottom * height_ratio)
    cv2.rectangle(image, (left, top), (right, bottom), class_colors[label], 2)
    cv2.putText(image, "{} {:.2f}".format(label, float(confidence)),
                (left, top - 5), cv2.FONT_HERSHEY_SIMPLEX, 0.5,
                class_colors[label], 2)
cv2.imshow(image)
```

```
# function to convert the JavaScript object into an OpenCV image
def js_to_image(js_reply):
    """
    Params:
        js_reply: JavaScript object containing image from webcam
    Returns:
        img: OpenCV BGR image
    """
    # decode base64 image
    image_bytes = b64decode(js_reply.split(',')[1])
    # convert bytes to numpy array
    jpg_as_np = np.frombuffer(image_bytes, dtype=np.uint8)
    # decode numpy array into OpenCV BGR image
    img = cv2.imdecode(jpg_as_np, flags=1)

    return img

# function to convert OpenCV Rectangle bounding box image into base64 byte string to be overlayed on video stream
def bbox_to_bytes(bbox_array):
    """
    Params:
        bbox_array: Numpy array (pixels) containing rectangle to overlay on video stream.
```

```

Returns:
    bytes: Base64 image byte string
"""
# convert array into PIL image
bbox_PIL = PIL.Image.fromarray(bbox_array, 'RGBA')
iobuf = io.BytesIO()
# format bbox into png for return
bbox_PIL.save(iobuf, format='png')
# format return string
bbox_bytes = 'data:image/png;base64,{}'.format(str(b64encode(iobuf.getvalue()), 'utf-8'))

return bbox_bytes

```

JavaScript to properly create our live video stream using our webcam as input

```

def video_stream():
    js = Javascript('''
        var video;
        var div = null;
        var stream;
        var captureCanvas;
        var imgElement;
        var labelElement;

        var pendingResolve = null;
        var shutdown = false;

        function removeDom() {
            stream.getVideoTracks()[0].stop();
            video.remove();
            div.remove();
            video = null;
            div = null;
            stream = null;
            imgElement = null;
            captureCanvas = null;
            labelElement = null;
        }

        function onAnimationFrame() {
            if (!shutdown) {
                window.requestAnimationFrame(onAnimationFrame);
            }
            if (pendingResolve) {
                var result = "";
                if (!shutdown) {
                    captureCanvas.getContext('2d').drawImage(video, 0, 0, 640, 480);
                    result = captureCanvas.toDataURL('image/jpeg', 0.8)
                }
                var lp = pendingResolve;
                pendingResolve = null;
                lp(result);
            }
        }

        async function createDom() {
            if (div !== null) {
                return stream;
            }

            div = document.createElement('div');
            div.style.border = '2px solid black';
            div.style.padding = '3px';
            div.style.width = '100%';
            div.style.maxWidth = '600px';
            document.body.appendChild(div);

            const modelOut = document.createElement('div');
            modelOut.innerHTML = "<span>Status:</span>";
            labelElement = document.createElement('span');
            labelElement.innerText = 'No data';
            labelElement.style.fontWeight = 'bold';
            modelOut.appendChild(labelElement);
            div.appendChild(modelOut);

            video = document.createElement('video');
            video.style.display = 'block';
            video.width = div.clientWidth - 6;
            video.setAttribute('playsinline', '');
            video.onclick = () => { shutdown = true; };
            stream = await navigator.mediaDevices.getUserMedia(
                {video: { facingMode: "environment"}});
            div.appendChild(video);

```



```

imgElement = document.createElement('img');
imgElement.style.position = 'absolute';
imgElement.style.zIndex = 1;
imgElement.onclick = () => { shutdown = true; };
div.appendChild(imgElement);

const instruction = document.createElement('div');
instruction.innerHTML =
  '<span style="color: red; font-weight: bold;">' +
  'When finished, click here or on the video to stop this demo</span>';
div.appendChild(instruction);
instruction.onclick = () => { shutdown = true; };

video.srcObject = stream;
await video.play();

captureCanvas = document.createElement('canvas');
captureCanvas.width = 640; //video.videoWidth;
captureCanvas.height = 480; //video.videoHeight;
window.requestAnimationFrame(onAnimationFrame);

return stream;
}
async function stream_frame(label, imgData) {
  if (shutdown) {
    removeDom();
    shutdown = false;
    return '';
  }

  var preCreate = Date.now();
  stream = await createDom();

  var preShow = Date.now();
  if (label != "") {
    labelElement.innerHTML = label;
  }

  if (imgData != "") {
    var videoRect = video.getClientRects()[0];
    imgElement.style.top = videoRect.top + "px";
    imgElement.style.left = videoRect.left + "px";
    imgElement.style.width = videoRect.width + "px";
    imgElement.style.height = videoRect.height + "px";
    imgElement.src = imgData;
  }

  var preCapture = Date.now();
  var result = await new Promise(function(resolve, reject) {
    pendingResolve = resolve;
  });
  shutdown = false;

  return {'create': preShow - preCreate,
    'show': preCapture - preShow,
    'capture': Date.now() - preCapture,
    'img': result};
}
'''

display(js)

def video_frame(label, bbox):
  data = eval_js('stream_frame("{}","{}").format(label, bbox)')
  return data

```

```

# start streaming video from webcam
video_stream()
# label for video
label_html = 'Capturing...'
# initialize bounding box to empty
bbox = ''
count = 0
while True:
  js_reply = video_frame(label_html, bbox)
  if not js_reply:
    break

  # convert JS response to OpenCV Image
  frame = js_to_image(js_reply["img"])

```

```
# create transparent overlay for bounding box
bbox_array = np.zeros([480,640,4], dtype=np.uint8)

# call our darknet helper on video frame
detections, width_ratio, height_ratio = darknet_helper(frame, width, height)

# loop through detections and draw them on transparent overlay image
for label, confidence, bbox in detections:
    left, top, right, bottom = bbox2points(bbox)
    left, top, right, bottom = int(left * width_ratio), int(top * height_ratio), int(right * width_ratio), int(bottom * height_ratio)
    bbox_array = cv2.rectangle(bbox_array, (left, top), (right, bottom), class_colors[label], 2)
    bbox_array = cv2.putText(bbox_array, "{} {:.2f}".format(label, float(confidence)),
                            (left, top - 5), cv2.FONT_HERSHEY_SIMPLEX, 0.5,
                            class_colors[label], 2)

bbox_array[:, :, 3] = (bbox_array.max(axis = 2) > 0 ).astype(int) * 255
# convert overlay of bbox into bytes
bbox_bytes = bbox_to_bytes(bbox_array)
# update bbox so next frame gets new overlay
bbox = bbox_bytes
```